

# Gaussian Process Regression in Finance: From Dynamic Incentive Models to Portfolio Optimization

Simon Scheidegger<sup>1</sup>

<sup>1</sup>Department of Economics, University of Lausanne

Deep Learning for Economics and Finance  
University of Geneva; March 25th - 27th, 2025

# Outline of the Talk

1. Introduction and Motivation
2. Gaussian Process Regression: A Primer
3. Two Dynamic Models
  1. Machine Learning for Dynamic Incentive Problems (Renner and Scheidegger, 2018)
  2. A Comprehensive Machine Learning Framework for Dynamic Portfolio Choice With Transaction Costs (Gaegauf et al., 2023)
4. Key Takeaways and Conclusion

# Dynamic Models in Finance and Economics

Many contemporary research questions are addressed via dynamic structural models, such as:

- ▶ Dynamic Incentive Problems
- ▶ Monetary Policy
- ▶ Climate Change

→ **Rich formulated economic environments:**

- ▶ Models are often **stochastic**
- ▶ Models consist of **many agents**
- ▶ **Nonlinear** (e.g., due to financial frictions)
- ▶ **Many Parameters** (structural estimation)



# Our Friend, the Bellman Equation

- **Dynamic Models** are often defined via the Bellman equation, characterizing the value function  $V(s)$  for a state  $s$ .


$$V(s) = \max_{a \in A(s)} \left\{ R(s, a) + \beta \int_S V(s') P(s'|s, a) ds' \right\}$$

- $V(s)$ : Value function, representing maximum expected reward from state  $s$ .
- $a$ : Action taken in state  $s$ .
- $A(s)$ : Set of allowable actions in  $s$ .
- $R(s, a)$ : Immediate reward for action  $a$  in state  $s$ .
- $\beta$ : Discount factor,  $0 < \beta \leq 1$ , weighting future rewards.
- $P(s'|s, a)$ : Transition probability to state  $s'$  given  $s$  and  $a$ .
- $S$ : State space, which can be  $d$ -dimensional with large  $d$ .

→ **This recursive equation is often solved with Value Function Iteration.**

# Numerical Roadblocks of Dynamic Models

1. Models suffer from the curse of dimensionality (many state variables,...).
2. Models suffer from non-linearities (financial frictions,...).
3. Have to approximate and interpolate high-dimensional functions on irregular-shaped geometries.
4. For high dimensions:  $\text{ratio of Volume(Sphere)/Volume(Cube)} \rightarrow 0$ .
5. **If projection methods/DP are used, solving optimization problems is expensive.**



→ **One way out: Gaussian Processes** (Rasmussen and Williams, 2005).


# Gaussian Process Regression: A Primer

# Function Approximation with Gaussian Processes

- ▶ Want to approximate Value- and Policy functions
  - ▶ in high dimensions
  - ▶ that are nonlinear
  - ▶ on irregular geometries
- ▶ Gaussian Process (GP) regression is a non-parametric method used to approximate functions based on observed data that can satisfy this.
- ▶ Given training data  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ , we aim to infer the underlying function  $f(x)$ .

$$y_i = f(x_i) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2)$$

# Intuition for GPs



We assume that  $f$ 's (heights) are **Gaussian distributed** with zero mean and some covariance matrix  $K$ :

$$\begin{pmatrix} f_1 \\ f_2 \\ f_3 \end{pmatrix} \sim \mathcal{N} \left( \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}, \begin{bmatrix} K_{11} & K_{12} & K_{13} \\ K_{21} & K_{22} & K_{23} \\ K_{31} & K_{32} & K_{33} \end{bmatrix} \right)$$

Note:  $f_1$  and  $f_2$  should be more correlated due to proximity (compared to  $f_1$  and  $f_3$ ). Example:

$$\sim \mathcal{N} \left( \begin{pmatrix} 0 \\ 0 \\ 0 \end{pmatrix}, \begin{bmatrix} 1 & 0.7 & 0.2 \\ 0.7 & 1 & 0.6 \\ 0.2 & 0.6 & 1 \end{bmatrix} \right)$$

Covariance matrix based on a **kernel function**, e.g.,:

$$\kappa(x, x') = \sigma_f^2 \exp \left( -\frac{1}{2\ell^2} (x - x')^2 \right)$$

where:

- ▶  $\sigma_f^2$ : Controls vertical variation.
- ▶  $\ell$ : Controls horizontal length scale.



# Defining the Gaussian Process Prior

- ▶ A GP is defined by a **mean function**  $m(x)$  and a **covariance function**  $k(x, x')$ :

$$f(x) \sim \mathcal{GP}(m(x), k(x, x'))$$

- ▶ Commonly, we set  $m(x) = 0$  to simplify computation, focusing on the covariance function.
- ▶ For any set of inputs  $\{x_1, \dots, x_N\}$ , **the function values**  $\{f(x_1), \dots, f(x_N)\}$  follow a **multivariate Gaussian distribution**:

$$\mathbf{f} \sim \mathcal{N}(0, K)$$

where  $K_{ij} = k(x_i, x_j)$ .

# Posterior Inference for Gaussian Process Regression

- ▶ Given training data  $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$ , we want to predict  $f(x^*)$  at a new point  $x^*$ .
- ▶ The joint distribution of observed values  $\mathbf{y}$  and the prediction  $f(x^*)$  is:

$$\begin{bmatrix} \mathbf{y} \\ f(x^*) \end{bmatrix} \sim \mathcal{N} \left( 0, \begin{bmatrix} K + \sigma^2 I & k_* \\ k_*^\top & k(x^*, x^*) \end{bmatrix} \right)$$

where  $k_* = [k(x^*, x_1), \dots, k(x^*, x_N)]$ .

- ▶ The conditional distribution of  $f(x^*) \mid \mathbf{y}$  is:

$$f(x^*) \mid \mathbf{y} \sim \mathcal{N}(\mu_{f(x^*)}, \sigma_{f(x^*)}^2),$$

where the predictive mean is given by:

$$\mu_{f(x^*)} = k_*^\top (K + \sigma^2 I)^{-1} \mathbf{y}$$

and the predictive variance reads as:

$$\sigma_{f(x^*)}^2 = k(x^*, x^*) - k_*^\top (K + \sigma^2 I)^{-1} k_*$$

# Predictive Mean and Variance


- To interpolate the function at any new point  $x^*$ , we use the **\*\*predictive mean\*\*** of the GP as our **best estimate**:

$$\text{Predictive Mean: } \mathbb{E}[f(x^*)] = k_*^\top (K + \sigma^2 I)^{-1} \mathbf{y}.$$

- The **\*\*predictive variance\*\*** quantifies the uncertainty around this estimate, providing **confidence intervals**:

$$\text{Predictive Variance: } \text{Var}(f(x^*)) = k(x^*, x^*) - k_*^\top (K + \sigma^2 I)^{-1} k_*.$$

- **Together, these expressions allow us to both estimate the function's value at any  $x^*$  and understand the associated uncertainty.**



# Choosing the Prior Kernel Function

- ▶ The kernel function  $k(x, x')$  determines the properties of the GP, such as smoothness and periodicity.
- ▶ Common choices include:
  1. **Squared Exponential (SE) Kernel:**

$$k(x, x') = \sigma^2 \exp \left( -\frac{(x - x')^2}{2\ell^2} \right)$$

where  $\ell$  controls smoothness and  $\sigma$  controls output variance.

2. **Matérn Kernel** (provides more flexibility in smoothness):

$$k(x, x') = \sigma^2 \frac{2^{1-\nu}}{\Gamma(\nu)} \left( \frac{\sqrt{2\nu}|x - x'|}{\ell} \right)^\nu K_\nu \left( \frac{\sqrt{2\nu}|x - x'|}{\ell} \right)$$

where  $\nu$  determines the differentiability.

3. **Periodic Kernel:** captures repeating patterns, useful in seasonal data.
4. Kernel Cookbook: <https://www.cs.toronto.edu/~duvenaud/cookbook/>

# Hyperparameter Optimization in GPR

- ▶ **Goal:** Optimize the hyperparameters of the kernel function to improve the fit of the GP model.
- ▶ **Process:** Hyperparameters are tuned by maximizing the marginal likelihood of the observed data:

$$\log p(\mathbf{y}|\mathbf{X}, \theta) = -\frac{1}{2}\mathbf{y}^\top (K_\theta + \sigma^2 I)^{-1} \mathbf{y} - \frac{1}{2} \log \det(K_\theta + \sigma^2 I) - \frac{n}{2} \log(2\pi)$$

where  $\theta$  represents the hyperparameters of the kernel  $K_\theta$ .

- ▶ **Common Optimizers:**
  - ▶ **Gradient Descent:** Uses the gradient of the marginal likelihood to adjust hyperparameters iteratively.
  - ▶ **L-BFGS (Limited-memory Broyden-Fletcher-Goldfarb-Shanno):** A popular quasi-Newton method efficient for high-dimensional optimization.
  - ▶ **Adam (Adaptive Moment Estimation):** An adaptive learning rate optimizer often used in machine learning, balancing speed and stability.

# Bayesian Active Learning with Gaussian Processes

→ **Uniform sampling in a high-dimensional space to populate the training set might be inefficient.**

- ▶ **Goal:** Given a fixed computational budget, efficiently approximate functions using GPs by **actively selecting data points that improve the model's accuracy.**
- ▶ **Core Idea:** Rather than uniformly sampling, **BAL strategically adds observations in areas where they contribute most to the model's performance** (e.g., MacKay, 1992, Chaloner and Verdinelli, 1995, Krause et al., 2008, Deisenroth et al., 2009, and Makarova et al., 2022).
- ▶ **Application:** Particularly useful in settings like value function iteration, where acquiring **new data points is computationally costly** (each training sample consists of solving a constrained optimization problem).

# How Bayesian Active Learning Works

1. **Model Update:** Fit a GP to current data  $\mathcal{D} = \{\mathbf{X}, \mathbf{y}\}$  to obtain a posterior mean and uncertainty for each point.
2. **Acquisition Function:** Evaluate candidate points by their expected contribution to the model, often using a score like:


$$\alpha(\tilde{x}_j; \mathcal{D}) = \sigma_m \tilde{\mu}(\tilde{x}_j) + \frac{\sigma_v}{2} \log(\tilde{\sigma}(\tilde{x}_j))$$

3. **Data Selection:** Select the candidate with the highest score, add it to the dataset, and update the model.<sup>1</sup>

---

<sup>1</sup>For effective use of the acquisition function, one needs to \*\*randomly sample a set of candidate points\*\* (e.g., 100), rank them by their acquisition scores, and add the highest-ranked point(s) to the training set.


# BAL Process Illustrated



**Figure:** Phased visualization of Bayesian Active Learning.




# BAL Process Illustrated




**Figure:** Phased visualization of Bayesian Active Learning.

# BAL Process Illustrated



**Figure:** Phased visualization of Bayesian Active Learning.

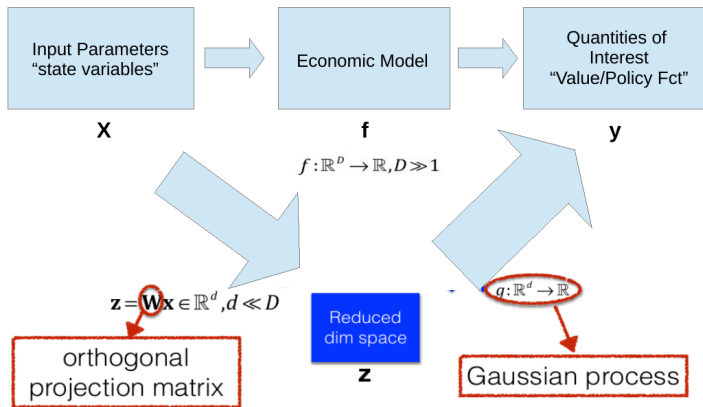
# Summary of Bayesian Active Learning Process



**Figure:** Phased Bayesian Active Learning: Initial data, refined with incremental observations via BAL.

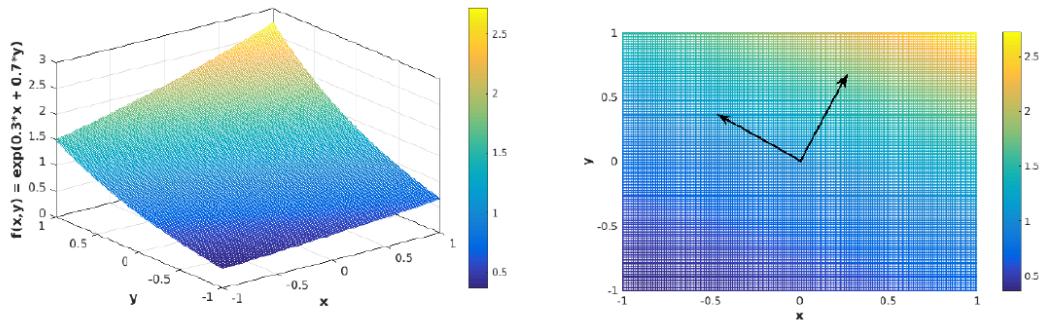
# Active Subspace-Based Dimension Reduction

- ▶ For dimensions larger than about 10, GPs drastically lose performance (Euclidean distance becomes uninformative in high-dim spaces).
- ▶ Need a formal way for reducing dimensionality.
- ▶ See Constantine et al. (2014) in general, Scheidegger and Bilonis (2019) for DP.



**Figure:** Intuition behind Active Subspaces.

# Active Subspaces – intuitive example



**Figure:** Function varies most along  $[0.3, 0.7]$ , and is constant in the orthogonal direction.

# Discover Active Subspaces

See, e.g., Constantine (2015), with references therein

## Step 1: Find $\mathbf{W}$

$$\mathbf{C} = \mathbb{E} \left[ \nabla_x f(x) \nabla_x f(x)^T \right] \approx \frac{1}{N} \sum_{i=1}^N \nabla_x f(x^{(i)}) \nabla_x f(x^{(i)})^T$$

$$\mathbf{C} = \mathbf{W} \mathbf{\Lambda} \mathbf{W}^T$$

- ▶ “Mean-square directional derivative”
- ▶ Note: derivative-free methods also exist for constructing  $\mathbf{W}$ .

# Discover Active Subspaces (cont.)

## Step 2: Partition the Eigendecomposition

$$\Lambda = \begin{bmatrix} \Lambda_1 & \\ & \Lambda_2 \end{bmatrix}, \quad \mathbf{W} = [\mathbf{W}_1 \quad \mathbf{W}_2], \quad \mathbf{W}_1 \in \mathbb{R}^{m \times n}$$

- ▶ Compute Eigenvalues, order them.
- ▶ Look for “gaps.”


## Step 3: Create a Rotated Coordinate System

$$\mathbf{x} = \mathbf{W}\mathbf{W}^T \mathbf{x} = \mathbf{W}_1 \mathbf{W}_1^T \mathbf{x} + \mathbf{W}_2 \mathbf{W}_2^T \mathbf{x} = \mathbf{W}_1 \mathbf{q} + \cancel{\mathbf{W}_2 \mathbf{v}}$$

# Dimension Reduction: An Example

$$f : [-1, 1]^{10} \rightarrow \mathbb{R}$$

$$f(x_1, \dots, x_{10}) = \exp(0.01x_1 + \textcolor{red}{0.7x_2} + 0.02x_3 + 0.03x_4 + 0.04x_5 \\ + 0.05x_6 + 0.06x_7 + 0.08x_8 + 0.09x_9 + 0.1x_{10})$$



**Figure:** Active Subspaces in Action.



# Gaussian Processes in Perspective

- ▶ GPs, together with BAL (and active subspaces), provide a tool to approximate high-dimensional, nonlinear functions on irregular geometries.
- ▶ This is particularly useful when training data is costly to acquire (each data point corresponds to solving a constrained optimization problem).
- ▶ This combination provides a generic tool for value function iteration.

**Table:** GPs in perspective with other function approximators.

Properties	Gaussian Process Regression	Neural Networks	Adaptive Sparse Grids	Chebyshev Polynomials
Handles High Dimensionality	✓	✓	✓	×
Nonlinear Function Approximation	✓	✓	✓	✓
Probabilistic Predictions	✓	×	×	×
Scalability with Data Size	×	✓	✓	×
Requires Grid Construction	×	×	✓	×
Adaptive Refinement	✓	×	✓	×
Interpretability	✓	×	✓	✓
Parallelization Potential	✓	✓	✓	✓
Hyperparameter Tuning Needed	Moderate	High	Low	Low
Uncertainty Quantification	✓	×	Limited	×

# Two Dynamic Models

# 1. Dynamic Incentive Problems

## Dynamic incentive problems:

- ▶ They occur whenever **two parties with repeated interaction under asymmetric information form a contract**.
- ▶ Applications: insurance contracts, optimal taxation, manager remuneration, ...
- ▶ e.g., a **risk averse agent** buys insurance from a **risk neutral planner** against income shocks (which are hidden to the planner).
  - i) **adverse selection**: there is *hidden information* (to one party in the contract).
  - ii) **moral hazard**: there is a *hidden action*.

## Example below: dynamic adverse selection problems with persistent shocks

- Solving such models numerically is a formidable task.
- Models are often formulated in a *stylized* fashion to remain computationally **tractable**.

# A simple dynamic optimal insurance model

see e.g. *Fernandes and Phelan (2000); Golosov et al. (2016)*

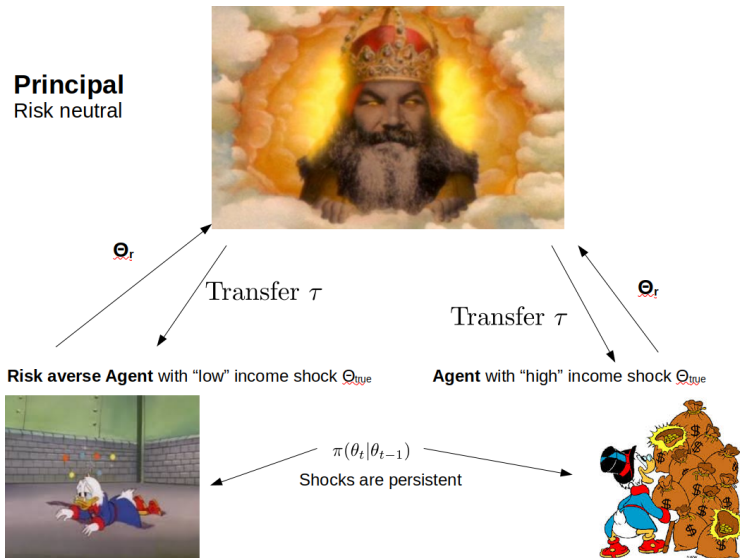
## A prototypical insurance model:

- ▶ Infinite-horizon, discrete time model.
- ▶ **Risk-averse agent:** faces *privately-observed* income shocks  $\Theta$ .
- ▶ Income shocks are **persistent**.
- ▶ **Risk neutral principal:** wants to provide optimal *incentive-compatible insurance* against income shocks.
- ▶ **Agent:** reports his income shock to the principal.
- ▶ **Principal:** *transfers consumption or charges a fee* to the agent, dependent on report.

## Optimal solution to the problem:

→ **A transfer scheme that maximizes the principal's utility while delivering a predetermined lifetime utility to the agent.**

# A Prototypical Dynamic Incentive Model



# An auxiliary problem formulation

## Formal issues:

- ▶ The observed reports are **fully history dependent**.
  - ▶ The principal needs to be prepared for infinitely many reporting strategies.
- **there is no obvious recursive formulation for the model.**

## To make the problem formally tractable (following existing literature):

- Apply the *revelation principle* and *one-shot deviation principle*.
- Introduce an auxiliary problem with **promised expected utilities** as state variables.
- To discipline the agent, we need **promise keeping** and **threat keeping** as states (lying always needs to be sub-optimal).
- For **every type of shock** in the model, we need an additional state (if there are  $N$  types, there are  $N - 1$  possibilities to lie).
- The continuous state space of the auxiliary problem is **N-dimensional**.

# Adverse selection: why difficult to solve?

Addressing **dynamic adverse selection models with persistent shocks** is a formidable task, since **two major computational bottlenecks** create difficulties in the solution process.

- 1) The **feasible sets** of utility promises **are not known in advance** and have to be **determined numerically**.
- 2) For solving dynamic adverse selection problems with value function iteration, we need to **repeatedly approximate and evaluate high-dimensional functions at arbitrary coordinates** within the feasible sets.

→ **Use-case for Gaussian Process Dynamic Programming.**

# Our solution method





- Generic penalty reformulation of dynamic adverse selection problems.

$$\begin{aligned}\bar{K}(\hat{v}(\cdot), \theta_-) &= \max_{c(\cdot), w(\cdot|\cdot), \xi_j} \sum_{\theta \in \Theta} \pi(\theta|\theta_-) [v(c(\theta), \theta_-) + \beta \bar{K}(w(\theta|\cdot), \theta)] - M \sum_{j=1}^m \xi_j^2 \\ \hat{v}(\theta_{(j)}) + \xi_j &= \sum_{\theta \in \Theta} \pi(\theta|\theta_{(j)}) [U(c(\theta), \theta) + \beta w(\theta|\theta)], \quad \forall j \in \{1, \dots, m\} \\ U(c(\theta), \theta) + \beta w(\theta|\theta) &\geq U(c(\hat{\theta}), \theta) + \beta w(\hat{\theta}|\theta), \quad \forall \theta, \hat{\theta} \in \Theta \\ c(\theta) \in C, w(\theta|\cdot) \in \mathcal{B} &= \prod_{\tilde{\theta} \in \Theta} [U(\underline{c}, \tilde{\theta})/(1-\beta), U(\bar{c}, \tilde{\theta})/(1-\beta)], \quad \forall \theta \in \Theta\end{aligned}$$

- This reformulation has the advantage that computing solutions to the reformulated problem reduces to solving an **ordinary dynamic programming problem via value function iteration**.
- **No need to explicitly characterize the feasible set.**
- **So far: Set-valued DP (Abreu et al., 1986); geometric constructions, don't scale.**
- Use GPs to approximate value and policy functions.





# Intuition of method (works for at least ten types)



## 2. Dynamic Portfolio Choice with Transaction Costs

- ▶ Optimization of intertemporal **utility subject** to dynamic **budget constraint**.
- ▶ Solution characterized by means of associated **Bellman equation**.
- ▶ Solvable in closed-form under **stringent assumptions** about, e.g., utility, return dynamics, market completeness, and cost structure.
- ▶ Existing numerical solution methods are limited by the **curse of dimensionality**.
- ▶ Existing numerical solution methods hardly applicable with more than **3-4 risky assets**.
- ▶ Intrinsically "**distinct**" optimal portfolio behavior over an **irregularly-shaped** domain.
- ▶ **Not everywhere differentiable** optimization problem with varying **smoothness** and **nonlinearity** properties over admissible domain.

# Some Issues: Curse of Dimensionality, and Irregularly-Shaped Domain



- ▶ Budget constraint depends on  $D$ -dimensional vector of risky assets' wealth shares.
- ▶ Irregularly-shaped domain with, e.g., no-short selling/borrowing constraints.
- ▶ A challenge for grid-based function approximation methods.

# Objective

A **comprehensive machine learning solution framework** for dynamic portfolio choice with proportional transaction costs (Gaegauf et al., 2023):

- ▶ **Scalable.**
  - ▶ Accomodating **nonlinearities**, **nondifferentiabilities** and **irregularly-shaped** domains.
  - ▶ "**Tunable-smooth**" over different subregions of the domain.
  - ▶ Allowing a **quantification** of the uncertainty of computed solutions.
  - ▶ **Extendable**, e.g., to general equilibrium and state-dependent opportunity sets.
- Use-case for GPs.

# Recursive form and Bellman equation

1. For some economically motivated **terminal utility** function  $U(\cdot, \cdot)$ , let:

$$V_T(W_T, \mathbf{x}_T) := U(W_T, \mathbf{x}_T) \quad (1)$$

2. Define **recursively**, for any  $t = T - 1, \dots, 0$ :

$$V_t(W_t, \mathbf{x}_t) = \max_{(C_t, \delta_t) \in \mathcal{D}(W_t, \mathbf{x}_t)} \{u(t, C_t) + \mathbb{E}[V(W_{t+1}, \mathbf{x}_{t+1}) | W_t, \mathbf{x}_t]\} , \quad (2)$$

where

$$(W_{t+1}, \mathbf{x}_{t+1}) = BC(W_t, \mathbf{x}_t, C_t, \delta_t; R_f, \mathbf{R}_{t+1})$$

3. Obtain desired value function and optimal controls from last **iteration** at  $t = 0$ .

# Key Issues for Computational Solution Framework

1. "Exact" computation of  $V_t(\cdot, \cdot)$  over entire domain  $\mathcal{D}_t$  is **infeasible/inappropriate**.
2. Need **approximation** of  $V_t(\cdot, \cdot)$  from suitable set of "exact" computations  $\{V_t(W_i, \mathbf{x}_i)\}_{i=1}^N$ :
  - ▶ **Error** in approximation of  $V_t(\cdot, \cdot)$  needs to be **quantifiable**.
  - ▶ **Accuracy** to bound error propagation in global solution as, e.g.,  $D$  or  $T$  grows.
  - ▶ Capture varying degrees of **nonlinearity/nondifferentiability** over domain  $\mathcal{D}_t$ .
  - ▶ Efficiency with "exact" computations performed only **within irregularly-shaped** domain  $\mathcal{D}_t$ .
3. **Grid-free probabilistic** approximation methods are preferable.

# Comprehensive Machine Learning Framework

- ▶ Based on recursive **grid-free probabilistic** approximation of value function  $V_t(\cdot, \cdot)$ .
- ▶ Two **building blocks**:

1. Accurate description of unknown **NTR**:

$$\Omega_t := \{(W_t, \mathbf{x}_t) \in \mathcal{D}_t : \delta_t^{opt} = \mathbf{0}\}$$

by means of suitable **approximate NTR**  $\hat{\Omega}_t$ .

2. Accurate approximation of  $V_t(\cdot, \cdot)|_{\hat{\Omega}_t}$  and  $V_t(\cdot, \cdot)|_{\mathcal{D}_t \setminus \hat{\Omega}_t}$  using two different associated **Gaussian Process Regressions (GPRs)**.

# Value Function Approximation with two GPRs

- ▶ Account for intrinsically **distinct** value function properties **inside/outside** NTR:

$$v_{1t} := v_t \mathbf{1}_{\hat{\Omega}_t} ; \quad v_{2t} := v_t \mathbf{1}_{\mathcal{D}_t \setminus \hat{\Omega}_t}$$

- ▶ "Disaggregated" Bellman equation: For any  $t = T - 1, \dots, 0$ :

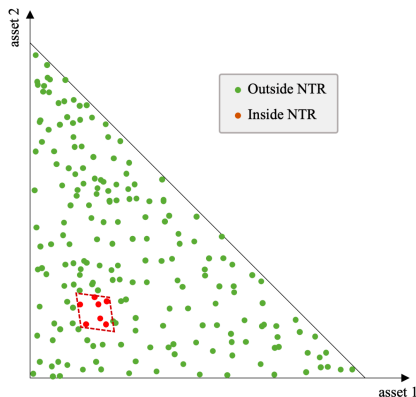
$$v_{1t}(\mathbf{x}_t) + v_{2t}(\mathbf{x}_t) = \max_{(c_t, \delta_t) \in \mathcal{D}(\mathbf{x}_t)} \left\{ u(c_t) + \beta \mathbb{E} \left[ \pi_{t+1}^{1-\gamma} (v_{1t+1}(\mathbf{x}_{t+1}) + v_{2t+1}(\mathbf{x}_{t+1})) \mid \mathbf{x}_t \right] \right\}$$

- ▶ Two distinct GPRs for approximating the value function **inside/outside** NTR:

$$\mathcal{GP}_{1t}(\mathbf{x}_t) + \mathcal{GP}_{2t}(\mathbf{x}_t) = \max_{(c_t, \delta_t) \in \mathcal{D}(\mathbf{x}_t)} \left\{ u(c_t) + \beta \mathbb{E} \left[ \pi_{t+1}^{1-\gamma} (\mathcal{GP}_{1t+1}(\mathbf{x}_{t+1}) + \mathcal{GP}_{2t+1}(\mathbf{x}_{t+1})) \right] \right\}$$



# Putting Things Together



- ▶ Randomly generate **training samples** of points inside/outside NTR.
- ▶ Train a **distinct GPR** on each subsample and **predict** remaining value function and policy values inside/outside NTR.

# Algorithm

**Input:** Terminal value function  $v_T$ , time horizon  $T$ , sampling size  $N$

**Output:** Surrogate value functions  $\{\mathcal{V}_{t-1}\}_{t=1}^T$  and approximate NTRs  $\{\hat{\Omega}_{t-1}\}_{t=1}^T$

1. Set  $\mathcal{V}_T = v_T$

2. **For**  $t = T$  **to** 1 **do**:

*Compute approximate NTR  $\hat{\Omega}_{t-1}$  using  $\mathcal{V}_t$  as next period's value function and sample  $N$  points  $\mathbf{X}_{t-1} = \{\mathbf{x}_{t-1i}\}_{i=1}^N \in \mathcal{D}_{t-1}^N$*

**For**  $i = 1$  **to**  $N$  **do**:

*Obtain value function and policy values  $(\hat{v}_{t-1i}, \hat{c}_{t-1i}, \hat{\delta}_{(t-1)i})$  for  $\mathbf{x}_{t-1i}$  by solving the Bellman equation using  $\mathcal{V}_t$  as next period's value function.*

**end**

Fill training sets  $\hat{\mathcal{D}}_{1t}, \hat{\mathcal{D}}_{2t}$  based on whether  $\mathbf{x}_{t-1i} \in \hat{\Omega}_{t-1}$  or  $\mathbf{x}_{t-1i} \notin \hat{\Omega}_{t-1}$  for some  $i$

Learn from  $\hat{\mathcal{D}}_{1t}, \hat{\mathcal{D}}_{2t}$  two GPR surrogates  $\mathcal{V}_{1(t-1)}, \mathcal{V}_{2(t-1)}$  of  $v_{1(t-1)}, v_{2(t-1)}$ , to finally obtain a surrogate  $\mathcal{V}_{t-1} = \mathcal{V}_{1(t-1)} + \mathcal{V}_{2(t-1)}$  of  $v_{t-1}$ .

**end**

# Key Takeaways and Conclusion

# Main Points to Remember

## Common Methodological Takeaways:

- ▶ Combining **dynamic programming** with **GPs** for efficient approximation of value and policy functions on irregular geometries is a powerful, scalable tool for solving dynamic models when data is expensive to acquire.
- ▶ **BAL** can boost performance by strategically focusing computational resources.
- ▶ **Active Subspaces** can help to reduce the dimensionality of problems effectively.
- ▶ **Grid-free approach** provides flexibility.
- ▶ Allows the **characterization of irregular state spaces**, such as no-trade regions or feasible sets in high-dimensional settings.

## Examples Highlighting the Method's Versatility:

- ▶ **Dynamic Incentive Problems (Renner and Scheidegger, 2018)**: Presents a solution framework for dynamic adverse selection models with **persistent types**, offering, e.g., insights into adverse selection effects in insurance markets.
- ▶ **Dynamic Portfolio Choice (Gaegauf et al., 2023)**: Applies the framework to **multi-asset portfolio optimization with transaction costs**, demonstrating that a broader asset space can mitigate liquidity constraints due to transaction costs.

Thank you! Questions?

# References I

- Abreu, D., Pearce, D., and Stacchetti, E. (1986). Optimal cartel equilibria with imperfect monitoring. *Journal of Economic Theory*, 39(1):251 – 269.
- Chaloner, K. and Verdinelli, I. (1995). Bayesian experimental design: A review. *Statist. Sci.*, 10(3):273–304.
- Constantine, P. G., Dow, E., and Wang, Q. Q. (2014). Active subspace methods in theory and practice: Applications to kriging surfaces (vol 36, pg a1500, 2014). *Siam Journal on Scientific Computing*, 36(6):A3030–A3031.
- Deisenroth, M. P., Rasmussen, C. E., and Peters, J. (2009). Gaussian process dynamic programming. *Neurocomputing*, 72(7):1508–1524.
- Fernandes, A. and Phelan, C. (2000). A recursive formulation for repeated agency with history dependence. *Journal of Economic Theory*, 91(2):223 – 247.
- Gaegauf, L., Scheidegger, S., and Trojani, F. (2023). A comprehensive machine learning framework for dynamic portfolio choice with transaction costs. *Available at SSRN 4543794*.

# References II

- Golosov, M., Tsyvinski, A., and Werquin, N. (2016). Recursive contracts and endogenously incomplete markets. *Handbook of Macroeconomics*, 2:725–841.
- Krause, A., Singh, A., and Guestrin, C. (2008). Near-optimal sensor placements in gaussian processes: Theory, efficient algorithms and empirical studies. *J. Mach. Learn. Res.*, 9:235–284.
- MacKay, D. J. C. (1992). Information-based objective functions for active data selection. *Neural Computation*, 4(4):590–604.
- Makarova, A., Shen, H., Perrone, V., Klein, A., Faddoul, J. B., Krause, A., Seeger, M., and Archambeau, C. (2022). Automatic termination for hyperparameter optimization. Technical report.
- Rasmussen, C. E. and Williams, C. K. I. (2005). *Gaussian Processes for Machine Learning (Adaptive Computation and Machine Learning)*. The MIT Press.
- Renner, P. and Scheidegger, S. (2018). Machine learning for dynamic incentive problems. Working paper. Available at SSRN: <http://dx.doi.org/10.2139/ssrn.3282487>.

# References III

Scheidegger, S. and Bilonis, I. (2019). Machine learning for high-dimensional dynamic stochastic economies. *Journal of Computational Science*, 33:68 – 82.