

A Tutorial on Physics-Informed Neural Networks (PINNs)

From 1D ODEs to 2D PDEs, Black–Scholes, and HJB

Simon Scheidegger

Department of Economics, University of Lausanne

Thursday, March 26th, 2025

Table of Contents

- 1 Introduction to PINNs
- 2 1D ODE: Zero Boundary Conditions
- 3 1D ODE: Non-Zero Boundary Conditions
- 4 2D PDE with Zero Boundary Conditions
- 5 2D PDE with Non-Zero Boundary Conditions
- 6 The Black–Scholes PDE
- 7 Hamilton–Jacobi–Bellman (HJB) Equation
- 8 Conclusion

What Are Physics-Informed Neural Networks (PINNs)?

- PINNs incorporate physical laws (ODEs, PDEs, etc.) into the training of neural networks.
- Instead of (or in addition to) fitting standard labeled data, we penalize the network if it violates the governing equations.
- Use automatic differentiation (AD) to compute derivatives for the PDE or ODE residual.
- Boundary/initial conditions are enforced as part of the loss function.
- We effectively transform solving differential equations into a data-driven, neural-network–based learning problem.

General PINN Strategy

- 1 Define a neural network $u_\theta(x)$ (or $u_\theta(x, y)$, etc. for PDEs).
- 2 Use automatic differentiation to compute partial derivatives needed for the PDE/ODE.
- 3 Define the **residual** for the equation, e.g. $F(u_\theta, x)$.
- 4 Incorporate boundary or initial conditions by adding penalty terms to the loss.
- 5 Train (optimize) to minimize:

$$\mathcal{L}(\theta) = \underbrace{\text{MSE}(F(u_\theta, x))}_{\text{Equation Residual}} + \underbrace{\text{MSE}(u_\theta(\text{boundary}) - \text{BC})}_{\text{Boundary Loss}} + \dots$$

- 6 After training, u_θ approximates the solution.

1D ODE with Zero Boundary Conditions

Example ODE:

$$\frac{d^2 y}{dx^2} = -1, \quad x \in (0, 1),$$

with **boundary conditions**

$$y(0) = 0, \quad y(1) = 0.$$

Analytical solution:

$$y(x) = -\frac{x^2}{2} + \frac{x}{2}.$$

PINN Setup:

- Define $y_\theta(x)$ as a neural net.
- PDE residual: $r(x) = y_\theta''(x) + 1$.
- Enforce boundary conditions $y_\theta(0) = 0$ and $y_\theta(1) = 0$.

Code Reference:

- `ode_zero_bc.py` demonstrates the full PyTorch implementation.
- It constructs a small network, sets up the ODE residual, and enforces boundary conditions within the loss.
- Training is run via Adam, and a final plot compares the learned solution to the analytical solution.

Slide Note: In live presentation, you can open or run `ode_zero_bc.py` to show details.

1D ODE with Non-Zero Boundary Conditions

Example ODE (same interior PDE, different BCs):

$$y''(x) = -1, \quad x \in (0, 1),$$

$$y(0) = 1, \quad y(1) = 2.$$

Analytical solution:

$$y(x) = -\frac{x^2}{2} + \frac{x}{2} + 1.$$

Main Difference:

- The boundary conditions now are $y(0) = 1$ and $y(1) = 2$.
- Only the boundary part of the loss changes.

Code Reference:

- `ode_nonzero_bc.py` shows the minor modifications in the boundary loss terms.

Example: 2D Laplace with Zero BCs

$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = 0, \quad (x, y) \in (0, 1) \times (0, 1),$$

with boundary $u = 0$ on $\partial\Omega$.

Trivial solution: $u(x, y) = 0$.

PINN approach:

- $u_\theta(x, y)$ as a neural net.
- Residual: $u_{xx} + u_{yy} = 0$.
- Enforce $u = 0$ on all edges.

Code Reference:

- `laplace_2d_zero_bc.py`.
- Implements interior sampling for (x, y) and boundary sampling on the edges.
- Laplace residual computed via second derivatives from 'torch.autograd.grad'.

Example: 2D Laplace with Non-Zero BCs

$$u_{xx} + u_{yy} = 0, \quad (x, y) \in (0, 1) \times (0, 1),$$

with $u(0, y) = 0$, $u(1, y) = 1$, $u(x, 0) = 0$, $u(x, 1) = 1$.

Analytical solution: $u(x, y) = \frac{x+y}{2}$.

Key difference:

- The boundary condition is no longer zero but $\{0, 1\}$.
- Adjust boundary losses accordingly.

Code Reference:

- `laplace_2d_nonzero_bc.py`.
- Shows how to define a boundary function for each edge and enforce it in the boundary part of the loss.

Black–Scholes Recap

$$\frac{\partial V}{\partial t} + \frac{1}{2}\sigma^2 S^2 \frac{\partial^2 V}{\partial S^2} + rS \frac{\partial V}{\partial S} - rV = 0.$$

$$\text{Terminal: } V(S, T) = \max(S - K, 0).$$

$$\text{Boundaries: } V(0, t) = 0, \quad V(S_{\max}, t) \approx S_{\max} - Ke^{-r(T-t)}.$$

Key Steps in the PINN Setup:

- PDE residual:

$$V_t + 0.5 \sigma^2 S^2 V_{SS} + rS V_S - r V.$$

- Boundary conditions at $S = 0$ and $S = S_{\max}$.
- Terminal condition at $t = T$.

Code Reference:

- `black_scholes_pinn.py` contains the full implementation.

We consider a 1D Hamilton–Jacobi–Bellman (HJB) PDE:

$$-r V(x) + \sup_{a \in [-1,1]} \left[(x + a) V'(x) - \alpha a^2 \right] = 0, \quad x \in (0, 1).$$

Boundary conditions:

$$V(0) = 0, \quad V(1) = 1.$$

Interpretation:

- $r > 0$ is a discount rate.
- $\alpha > 0$ penalizes the square of the control a .
- The state variable is $x \in [0, 1]$.

Optimal Control:

$$a^*(x) = \text{clamp}\left(\frac{V'(x)}{2\alpha}, -1, 1\right).$$

Then the PDE becomes:

$$-r V(x) + (x + a^*) V'(x) - \alpha (a^*)^2 = 0.$$

Key point: If $|V'(x)/(2\alpha)| \leq 1$, the unconstrained optimum is valid. Otherwise, it saturates at ± 1 .

- Let $V_\theta(x)$ be a fully connected neural network with parameters θ .
- We compute $V'_\theta(x)$ via **automatic differentiation**.
- PDE residual:

$$R_\theta(x) = -r V_\theta(x) + (x + a^*(x)) V'_\theta(x) - \alpha (a^*(x))^2.$$

- **Loss** for interior points:

$$\mathcal{L}_{\text{PDE}}(\theta) = \sum_{x_{\text{int}}} (R_\theta(x_{\text{int}}))^2.$$

- **Boundary conditions:**

$$V_\theta(0) = 0, \quad V_\theta(1) = 1 \quad \Rightarrow \quad \mathcal{L}_{\text{BC}}(\theta) = (V_\theta(0) - 0)^2 + (V_\theta(1) - 1)^2.$$

- **Total loss:** $\mathcal{L}(\theta) = \mathcal{L}_{\text{PDE}} + \mathcal{L}_{\text{BC}}$.

Implementation

- **Network architecture:** fully-connected, Tanh activations.
- **Optimization:** Adam with a moderate learning rate (e.g., 10^{-3}).
- **Sampling:**
 - **Interior points:** random $x \in (0, 1)$.
 - **Boundary points:** $x = 0$, $x = 1$.
- **Code Reference:**
 - See `hjb_pinn_continuous.py` for a complete example.
 - After training, we plot $V_\theta(x)$ on $x \in [0, 1]$.

- The PINN learns $V(x)$ that (approximately) satisfies the HJB PDE and boundary conditions.
- The learned policy is $a^*(x) = \text{clamp}(V'(x)/(2\alpha), -1, 1)$.
- No closed-form analytical solution is needed; we compare the training residual or do a quick check of boundary conditions.

Conclusion: PINNs can handle continuous controls in HJB equations by *differentiating* w.r.t. $V_\theta(x)$, computing the *optimal control*, and enforcing the PDE residual.

PDE Setup with non-trivial Boundary Conditions

We consider a 2D HJB:

$$-r V(x, y) + \max_{a \in [-1, 1]} \left[(x + a) \frac{\partial V}{\partial x} + (y + a) \frac{\partial V}{\partial y} - \alpha a^2 \right] = 0,$$

with $(x, y) \in (0, 1) \times (0, 1)$.

Non-trivial boundary conditions:

$$V(0, y) = y, \quad V(1, y) = 1 + y, \quad V(x, 0) = x, \quad V(x, 1) = x + 1.$$

- $r > 0$ is discount factor, $\alpha > 0$ is penalty on a^2 .
- PDE domain: the unit square.

Continuous Control in $[-1,1]$

Optimal control $a^*(x, y)$:

$$a_{\text{unclamped}} = \frac{\frac{\partial V}{\partial x} + \frac{\partial V}{\partial y}}{2\alpha}, \quad a^* = \text{clamp}(a_{\text{unclamped}}, -1, 1).$$

Then PDE residual:

$$R(x, y) = -r V + (x + a^*) \frac{\partial V}{\partial x} + (y + a^*) \frac{\partial V}{\partial y} - \alpha (a^*)^2 = 0.$$

Neural Network: $V_\theta(x, y)$.

- **Interior loss:**

$$\mathcal{L}_{\text{PDE}}(\theta) = \sum_{(x_i, y_i) \in \text{interior}} (R_\theta(x_i, y_i))^2.$$

- **Boundary loss:**

$$\mathcal{L}_{\text{BC}}(\theta) = \sum_{(x_j, y_j) \in \text{boundary}} \left(V_\theta(x_j, y_j) - \text{BC}(x_j, y_j) \right)^2,$$

where $\text{BC}(x, y)$ is the piecewise function that returns $y, 1 + y, x, x + 1$ on each edge.

- **Total loss:**

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{PDE}} + \mathcal{L}_{\text{BC}}.$$

Implementation

- **Network:** fully connected, e.g., Tanh activation, 64 hidden units.
- **Sampling:**
 - N_{interior} random points in $(0, 1) \times (0, 1)$.
 - N_{boundary} random points on each edge.
- **Optimizer:** Adam, with moderate LR (e.g. 10^{-3}).
- **Loss printing:** PDE loss, BC loss, total loss each iteration (or every 500 epochs).

Code Reference:

- See `hjb_2d_nontrivial_bc.py` for the full script.
- After training, we plot V_θ in a 2D color map and a 3D surface.

Summary 2D HJB with non-trivial BCs

- We introduced a 2D HJB PDE with non-trivial BCs.
- We clamp the control $a^* = \text{clamp}(\frac{V_x + V_y}{2\alpha}, -1, 1)$ inside the PDE.
- The PINN approach adds PDE residual loss + boundary mismatch loss.
- This yields $V_\theta(x, y)$ that approximates the PDE solution subject to the specified boundary condition.

Summary & Next Steps

- We introduced PINNs: enforcing PDE/ODE constraints by building them into the loss function.
- Showed ODEs, PDEs, boundary conditions, and terminal conditions.
- Demonstrated examples: 1D boundary value problems, 2D Laplace, Black–Scholes, and a simple HJB.
- In practice, you can adapt these templates to more complex PDEs, domains, or multi-dimensional states.

References & Further Reading

- M. Raissi, P. Perdikaris, and G. E. Karniadakis, Physics-Informed Neural Networks: A Deep Learning Framework (2019).
- J. Berg and K. Nystroem, A Unified Deep Artificial Neural Network Approach to PDEs in Complex Geometries (2017).
- For HJB PDE references in reinforcement-learning contexts, see: D. Jiang, F. Meng, Q. Sun, X. Xue, and Y. Zou. DeepRitz Method (2019).

Thank You!

Questions?