

# Recap: Physics-Informed Neural Networks

## Example: A 2D Poisson PDE, and Black-Scholes with 2 risky assets

Simon Scheidegger

March 27th, 2025

- **Objective:** Solve PDE by embedding the PDE residual into a neural network's loss.
- Instead of only fitting data, we also enforce:

$$\text{Loss} = \underbrace{\text{MSE}(\text{PDE residual})}_{\text{physics}} + \underbrace{\text{MSE}(\text{boundary conditions})}_{\text{boundary data}}.$$

- Use automatic differentiation (autograd) to compute PDE derivatives inside the network.
- Minimization via standard optimizers (e.g. Adam) yields an approximate solution.

## 2D Poisson Problem:

We consider the domain  $\Omega = [0, 1] \times [0, 1]$  with the PDE

$$\Delta u(x, y) = 4, \quad (x, y) \in \Omega,$$

and the **Dirichlet boundary condition**

$$u(x, y) = x^2 + y^2, \quad (x, y) \in \partial\Omega.$$

## Analytical solution:

It is straightforward to check that

$$u^*(x, y) = x^2 + y^2$$

satisfies both the PDE

$$\Delta(x^2 + y^2) = 2 + 2 = 4$$

and the same boundary condition.

- **PDE Residual:**  $\Delta u - 4 = 0$
- **Boundary Mismatch:**  $u - (x^2 + y^2) = 0$

- **Network:** fully-connected, Tanh activations.
- **Inputs:**  $(x, y)$  coordinates.
- **Output:**  $u_\theta(x, y)$ , neural approximation.
- **Loss function:**

$$\mathcal{L}(\theta) = \frac{1}{N_{\text{int}}} \sum_{\text{int pts}} [\Delta u_\theta - 4]^2 + \frac{1}{N_{\text{bc}}} \sum_{\text{bc pts}} [u_\theta - (x^2 + y^2)]^2.$$

# Results & Accuracy

- We train with random points in the domain for PDE, and on edges for boundary conditions.
- **Metrics:** MSE of PDE residual (interior) and boundary mismatch.
- After training, we evaluate:

$$\text{MSE}_{\text{interior}} = \frac{1}{|\Omega_{\text{test}}|} \sum_{\Omega_{\text{test}}} [u_{\theta} - u^*]^2,$$

$$\text{MSE}_{\text{boundary}} = \frac{1}{|\partial\Omega_{\text{test}}|} \sum_{\partial\Omega_{\text{test}}} [u_{\theta} - u^*]^2.$$

- Typically, errors can reach  $10^{-5}$  to  $10^{-7}$  range, depending on network size and training steps.

# Basket Options: An Overview

- **Basket Options** are derivatives written on a weighted average of multiple assets.
- Consider a European call option on a basket:

$$\text{Payoff} = \max(w_1 S_1 + w_2 S_2 - K, 0)$$

- In our example, we choose:

$$w_1 = w_2 = 0.5, \quad K = 100, \quad T = 1.$$

- The assets have different volatilities and a nonzero correlation.

# Dynamics of Underlying Assets

Assume under the risk-neutral measure the asset prices follow Geometric Brownian Motion:

$$dS_1 = rS_1 dt + \sigma_1 S_1 dW_1,$$

$$dS_2 = rS_2 dt + \sigma_2 S_2 dW_2,$$

with

$$dW_1 dW_2 = \rho dt.$$

- $r$ : Risk-free rate.
- $\sigma_1, \sigma_2$ : Volatilities.
- $\rho$ : Correlation.

# Derivation of the PDE (I)

Using Itô's formula for a twice-differentiable function  $V(t, S_1, S_2)$ , we have

$$dV = V_t dt + V_{S_1} dS_1 + V_{S_2} dS_2 + \frac{1}{2} V_{S_1 S_1} (dS_1)^2 + \frac{1}{2} V_{S_2 S_2} (dS_2)^2 + V_{S_1 S_2} dS_1 dS_2.$$

Substitute the dynamics of  $S_1$  and  $S_2$ :

$$\begin{aligned} dV &= V_t dt + rS_1 V_{S_1} dt + rS_2 V_{S_2} dt \\ &\quad + \frac{1}{2} \sigma_1^2 S_1^2 V_{S_1 S_1} dt + \frac{1}{2} \sigma_2^2 S_2^2 V_{S_2 S_2} dt \\ &\quad + \rho \sigma_1 \sigma_2 S_1 S_2 V_{S_1 S_2} dt + \text{stochastic terms.} \end{aligned}$$

In a risk-neutral world, the discounted price  $e^{-rt} V(t, S_1, S_2)$  must be a martingale. Thus, the drift must vanish:

$$V_t + rS_1 V_{S_1} + rS_2 V_{S_2} + \frac{1}{2} \sigma_1^2 S_1^2 V_{S_1 S_1} + \frac{1}{2} \sigma_2^2 S_2^2 V_{S_2 S_2} + \rho \sigma_1 \sigma_2 S_1 S_2 V_{S_1 S_2} - rV = 0.$$



# Derivation of the PDE (II)

- The PDE is solved backwards in time, i.e.,  $t$  represents time-to-maturity.
- **Terminal condition:**

$$V(T, S_1, S_2) = \max(w_1 S_1 + w_2 S_2 - K, 0).$$

- For our example, with  $w_1 = w_2 = 0.5$  and  $K = 100$ , the terminal payoff is:

$$V(T, S_1, S_2) = \max(0.5S_1 + 0.5S_2 - 100, 0).$$

- To make the computational domain finite, we truncate  $S_1, S_2 \in [0, S_{\max}]$  (with, say,  $S_{\max} = 200$ ) and impose homogeneous Neumann conditions (zero flux) at the boundaries.

# Physics-Informed Neural Networks (PINNs)

- **Idea:** Approximate the solution  $V(t, S_1, S_2)$  with a neural network  $V_\theta(t, S_1, S_2)$ .
- **Loss Function:**

$$\begin{aligned}\mathcal{L}(\theta) = & \underbrace{\frac{1}{N_{\text{int}}} \sum_{i=1}^{N_{\text{int}}} \left| \mathcal{R}(t_i, S_{1,i}, S_{2,i}; \theta) \right|^2}_{\text{Interior (PDE) Loss}} \\ & + \underbrace{\frac{1}{N_{\text{term}}} \sum_{j=1}^{N_{\text{term}}} \left| V_\theta(T, S_{1,j}, S_{2,j}) - \text{Payoff}_j \right|^2}_{\text{Terminal Loss}} \\ & + \underbrace{\mathcal{L}_{\text{neum}}}_{\text{Neumann Loss}}.\end{aligned}$$

- **PDE Residual:**  $\mathcal{R} = V_t + \frac{1}{2}\sigma_1^2 S_1^2 V_{S_1 S_1} + \frac{1}{2}\sigma_2^2 S_2^2 V_{S_2 S_2} + \rho\sigma_1\sigma_2 S_1 S_2 V_{S_1 S_2} + rS_1 V_{S_1} + rS_2 V_{S_2} - rV$ .
- The training minimizes  $\mathcal{L}(\theta)$  using gradient-based optimization.

- **Step 1:** Define the neural network architecture  $V_{\theta}(t, S_1, S_2)$ .
- **Step 2:** Use automatic differentiation to compute the derivatives in the PDE residual.
- **Step 3:** Sample points from the interior and boundaries:
  - Interior points for the PDE residual.
  - Terminal points (at  $t = T$ ) for the payoff.
  - Spatial boundary points for enforcing Neumann conditions.
- **Step 4:** Construct the total loss as the sum of the interior, terminal, and Neumann losses.
- **Step 5:** Train the network by minimizing the total loss with an optimizer (e.g., Adam).

# Implementation Details: Code Overview

- **Network:** The class `PINN_BS` implements a feed-forward network with three inputs ( $t, S_1, S_2$ ) and one output  $V$ .
- **PDE Residual:** Function `bs_pde_residual` computes the residual using autograd.
- **Loss Functions:**
  - `interior_loss` samples interior points and evaluates the PDE residual loss.
  - `terminal_loss` enforces the terminal condition.
  - `neumann_loss` enforces homogeneous Neumann conditions on the spatial boundaries.
- **Training:** The function `train_pinn_bs` combines these losses and trains the model.
- **Visualization:** 3D scatter plots are generated at  $t = T$  to compare the exact terminal solution, the PINN prediction, and the absolute error.

# Results: Training Curves and Metrics

- During training, the interior loss, terminal loss, and Neumann loss are tracked.
- At convergence, we report:
  - $\text{MSE}_{\text{interior}}$  on the domain.
  - $\text{MSE}_{\text{terminal}}$  at  $t = T$ .
  - Neumann error on the spatial boundaries.
- These metrics indicate the overall accuracy of the PINN approximation.

# Visualization: 3D Scatter Plots at Terminal Time

- The terminal condition is  $V(T, S_1, S_2) = \max(0.5S_1 + 0.5S_2 - 100, 0)$ .
- We evaluate the PINN prediction on a grid in  $(S_1, S_2)$  at  $t = T$ .
- Three 3D scatter plots are produced:
  - 1 Exact terminal solution.
  - 2 PINN predicted terminal solution.
  - 3 Absolute error between prediction and exact.

# Conclusions and Future Work You could consider

- PINNs provide a meshfree approach to solve high-dimensional PDEs.
- The two-asset Black–Scholes PDE for basket options can be solved by embedding the PDE and boundary conditions into a loss function.
- The presented approach is verified against the known terminal payoff.
- Future work may include:
  - More sophisticated network architectures.
  - Adaptive sampling strategies.
  - Application to more complex multi-asset or American options.